

---

# OSGS

*Release 0.1*

**Tim Sutton**

**Sep 06, 2021**



**CONTENTS:**

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Installation</b>	<b>3</b>
<b>3</b>	<b>Indices and tables</b>	<b>9</b>
<b>4</b>	<b>Services</b>	<b>11</b>
<b>5</b>	<b>Indices and tables</b>	<b>29</b>
<b>6</b>	<b>Utilities</b>	<b>31</b>
<b>7</b>	<b>Indices and tables</b>	<b>37</b>
<b>8</b>	<b>Workflows</b>	<b>39</b>
<b>9</b>	<b>Indices and tables</b>	<b>41</b>
<b>10</b>	<b>OSM Mirror</b>	<b>43</b>
<b>11</b>	<b>Indices and tables</b>	<b>49</b>



## **INTRODUCTION**

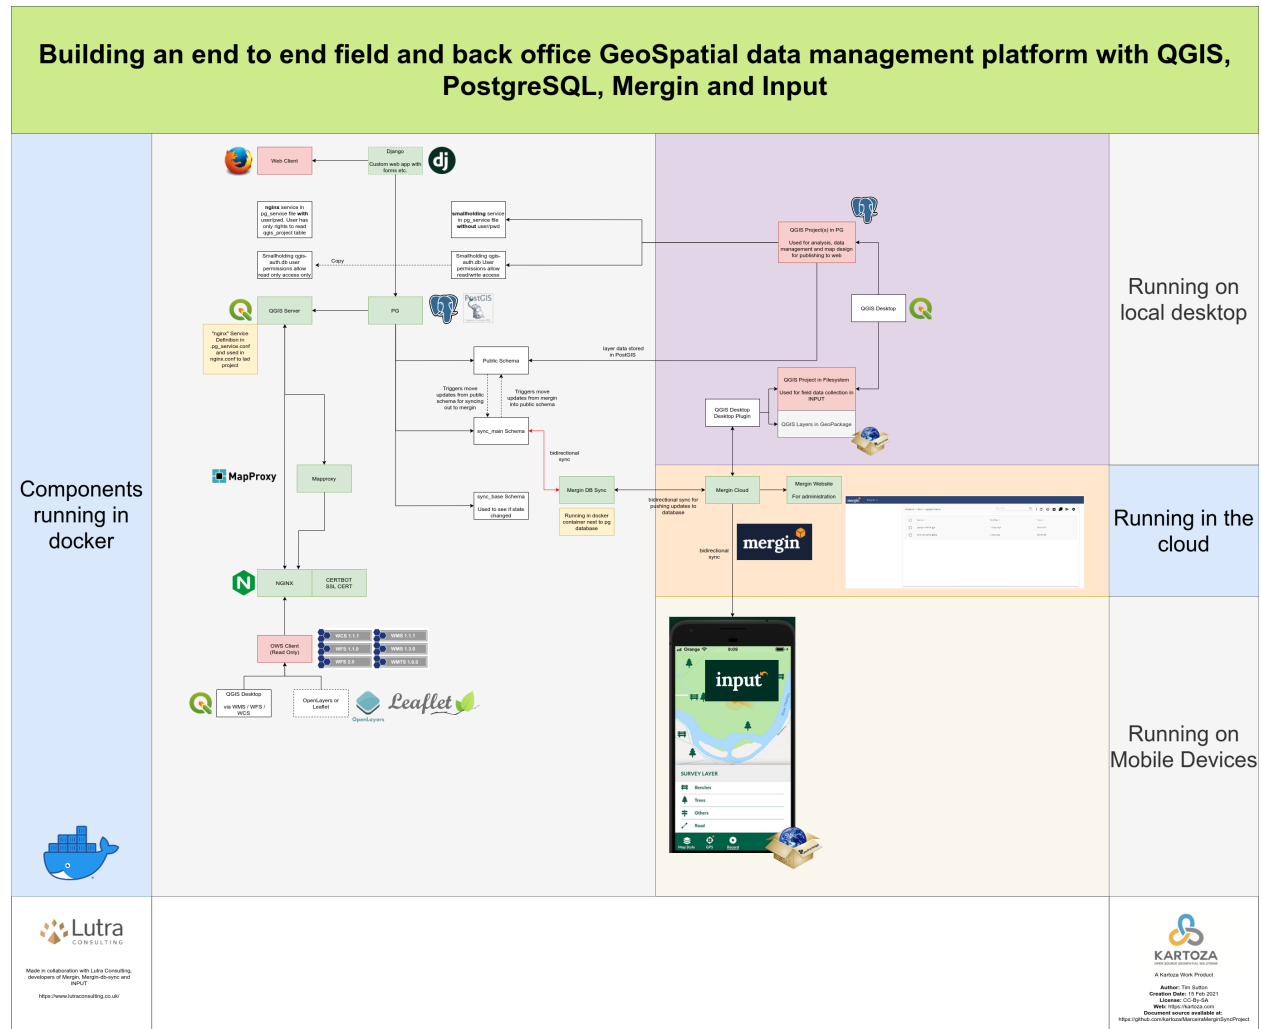
### **1.1 What is the Open Source GIS Stack?**

In the open source geospatial world, there are different interesting technologies that have been developed to allow one to manipulate, visualize, publish, manage, edit, etc. geospatial data. These technologies are often in disparate projects. Programs like the QGIS project, are often ensemble applications that take some of these different technologies as components and use them to provide some higher level functionality. In some cases the integration is in silos or vertical applications, for example QGIS, PostGIS, Geoserver, etc. The aim of this project is to take a number of those different silos and incorporate them into a single platform, therefore making it really easy to access and publish the different types of services that are easily available on their own but not so easily available in a consolidated platform.

The Open Source GIS Stack (OSGS) is a platform based on Docker, which is a container deployment environment. Containers are isolated application runtimes. They typically contain a minimal version of an operating system, typically Linux. In that minimal version of the operating system the containers will typically have one single task they are responsible for doing, for example running PostGIS or running Geoserver and so on. Each container has a very specific task and the design intent that you do not deploy multiple services in a single container. You must deploy each service in its own container then orchestrate those containers together. This is all managed using Docker and Docker Compose which is the orchestration environment for Docker. OSGS is simply a set of Docker orchestration routines to bring up different services in different containers, make them all talk together and present them as one cohesive architecture.

The other important thing that OSGS does, is it provides some scripts to do various management commands and tasks, for example, bringing up services or down services, copying data, making backups, etc. These scripts are provided using a utility called “make” which is typically used by programmers to compile software, but can also be used to automate tasks. There is a Makefile which does all the high level tasks on the orchestrated architecture that OSGS provides.

## 1.2 Overview Diagram



**Note:** Anybody can take open source software and package it as Docker services. Therefore, when you are choosing which Docker service and containers to run, you can go and have a look at the various ways different people have packaged up a particular software and find one that works the best for you.

## INSTALLATION

### 2.1 Preparing the server

#### 2.1.1 Basic Security

##### Unattended upgrades

This will automatically install only security fixes on a continual basis on your server.

```
sudo apt install unattended-upgrades
```

##### ssh

Disable password authentication for SSH

```
sudo vim /etc/ssh/sshd_config
```

Set this:

```
PasswordAuthentication no
Then do
sudo systemctl restart sshd.service
```

##### Crowdsec

<https://crowdsec.net/>

```
wget -q0 - https://s3-eu-west-1.amazonaws.com/crowdsec.debian.pragmatic/crowdsec.asc
↪|sudo apt-key add - && echo "deb https://s3-eu-west-1.amazonaws.com/crowdsec.debian.
↪pragmatic/$(lsb_release -cs) $(lsb_release -cs) main" | sudo tee /etc/apt/sources.list.
↪d/crowdsec.list > /dev/null;
sudo apt-get update
sudo apt-get install crowdsec
```

## Fail2ban

```
sudo apt install fail2ban
https://www.fail2ban.org/wiki/index.php/Main_Page
```

## Firewall

```
sudo ufw allow ssh
sudo ufw enable
sudo ufw status
```

Should show something like this:

```
Status: active

To                Action    From
--                -
22/tcp            ALLOW     Anywhere
22/tcp (v6)       ALLOW     Anywhere (v6)
```

We will open more ports as they are needed.

## Status monitoring

bpytop is a great console based dashboard for monitoring your server.

```
sudo snap install bpytop
```

## 2.1.2 Additional Software

### Docker

```
sudo apt install docker.io docker-compose
```

### Git, rpl, pwgen and Make. openssl

Needed for checking out our docker project and running the various make commands we provide.

```
sudo apt install git make rpl pwgen openssl
```

or fedora:

```
sudo dnf install openssl rpl git pwgen
```



## 2.2 Initial Configuration

Note for the unprivileged user throughout here, we use the user name ‘timlinux’ in various examples - you should substitute this with your own user.

### 2.2.1 User Group

Add yourself to the user group of docker so you don’t need to sudo docker commands.

```
sudo usermod -a -G docker timlinux
```

Then log out and in again to assume the upgraded permissions.

### 2.2.2 Project Checkout

```
cd /home
sudo mkdir web
sudo chown timlinux.timlinux web
cd web
git clone https://github.com/kartoza/OpenSource-GIS-Stack
cd OpenSource-GIS-Stack
```

### 2.2.3 Fetching Docker Images

```
[timlinux@fedora OpenSource-GIS-Stack]$ docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
apache/superset	latest-dev	dae66d8ef922	13 hours ago	1.9GB
pbf	stable	24a6ed68e359	14 hours ago	252MB
swaggerapi/swagger-ui	latest	be6494c7edf4	32 hours ago	98.7MB
lutraconsulting/mergin-db-sync	latest	2f479ae94f5d	2 days ago	630MB
opendronemap/odm	latest	223a72b0d4e5	3 days ago	1.62GB
byrnedo/alpine-curl	latest	5208e5295614	9 days ago	6.9MB
redis	latest	739b59b96069	10 days ago	105MB
nginx	latest	62d49f9bab67	2 weeks ago	133MB
node	14	d6602e31594f	2 weeks ago	943MB
kartoza/docker-osm	imposm-latest	3d8739408ed0	3 weeks ago	1.08GB
kartoza/docker-osm	osmupdate-latest	bb1457cbd311	3 weeks ago	250MB
certbot/certbot	latest	c8eac9ed295e	3 weeks ago	357MB
silexlabs/silex	latest	be2ad993c076	4 weeks ago	1.38GB
klakegg/hugo	0.82.0	2feec974b2db	5 weeks ago	46.1MB
openquake/qgis-server	stable	45d3c20d6088	2 months ago	1.26GB
kartoza/mapproxy	latest	d6df8e43db21	3 months ago	1.37GB
kartoza/postgis	13.0	a85a969d604b	5 months ago	1.34GB
kartoza/geoserver	2.18.0	f6774787e652	6 months ago	1.88GB
postgrest/postgrest	latest	befe6cd943da	11 months ago	84.6MB
kartoza/docker-osm	osmenrich-latest	ccab39593491	22 months ago	949MB
jguyomard/hugo-builder	latest	650ec6415cde	2 years ago	57.2MB
quay.io/lkiesow/docker-scp	latest	b99cf24fe937	2 years ago	12.5MB
bytemark/webdav	latest	c124350447bb	2 years ago	95.6MB
geodata/gdal	latest	1e1929f80f44	3 years ago	1.29GB

## 2.2.4 Configuration

Copy the .env boilerplate file and then adjust any settings in it as needed.

```
cp .env.example .env
```

Replace terms that should be unique for your installation:

```
rpl example.org geoservices.govt.lc .env  
rpl example.org geoservices.govt.lc nginx_conf/nginx.conf
```

## 2.3 Obtaining free fonts for your projects

There are two great sources of free fonts that you can use for your projects:

<http://ftp.gnu.org/gnu/freefont/freefont-ttf-20120503.zip> <https://fonts.google.com/>

There is a makefile target that can be run with:

```
make get-fonts
```

That will fetch all of these fonts for you. If you also fetch these fonts on your local machine and place them in your `~/ .fonts` folder then you can use them in your local QGIS projects, know they will also be available to QGIS server if you publish that project as a web map.

**Note:** This and other makefile targets assume that you have not changed the `COMPOSE_PROJECT_NAME=osgisstack` environment variable in .env.

**Note:** The above make command fetches a rather large download!

## 2.4 Production Stack

### 2.4.1 Overview

In this section we will bring up the full production stack, but to do that we first need to get an SSL certificate issued. To facilitate this, there is a special, simplified, version of Nginx which has no reverse proxies in place and not docker dependencies. Here is an overview of the process:

1. Replace the domain name in your letsencrypt init script
2. Replace the email address in your letsencrypt init script
3. Replace the domain name in the certbot init nginx config file
4. Open up ports 80 and 443 on your firewall
5. Run the init script, ensuring it completed successfully
6. Shut down the minimal nginx
7. Replace the domain name in the production nginx config file
8. Generate passwords for geoserver, postgres, postgres and update .env
9. Copy over the mapproxy template files
10. Run the production profile in docker compose

At the end of the process you should have a fully running production stack with these services:

IMAGE	PORTS	NAMES
kartoza/mapproxy	8080/tcp	osgisstack_mapproxy_1
nginx	0.0.0.0:80->80/tcp, 0.0.0.0:443->443/tcp	osgisstack_nginx_1
swaggerapi/swagger-ui	80/tcp, 8080/tcp	osgisstack_swagger_1
postgrest/postgrest	3000/tcp	osgisstack_postgrest_1
kartoza/geoserver:2.18.0	8080/tcp, 8443/tcp	osgisstack_geoserver_1
openquake/qgis-server:stable	80/tcp, 9993/tcp	osgisstack_qgis-server_1
kartoza/postgis:13.0	0.0.0.0:5432->5432/tcp	osgisstack_db_1
quay.io/lkiesow/docker-scp	0.0.0.0:2222->22/tcp	osgisstack_scp_1
certbot/certbot	80/tcp, 443/tcp	osgisstack_certbot_1

The following ports will be accessible on the host to the docker services. You can, on a case by case basis, allow these through your firewall using ufw (uncomplicated firewall) to make them publicly accessible:

1. 80 - http: Only really needed during initial setup of your letsencrypt certificate
2. 443 - https: All web based services run through this port so that they are encrypted
3. 5432 - postgres: Only expose this publicly if you intend to allow remote clients to access the postgres database.
4. 2222 - scp: There is an scp/sftp upload mechanism to mobilise data and resources to the web site

For those services that are not exposed to the host, they are generally made available over 443/SSL via reverse proxy in the Nginx configuration.

Some things should still be configured manually and deployed after the initial deployment:

1. Mapproxy configuration
2. setup.sql (especially needed if you are planning to use postgrest)
3. Hugo content management
4. Landing page static HTML

And some services are not intended to be used as long running services. especially the ODM related services.

## 2.4.2 Configuration

We have written a make target that automates steps 1-10 described in the overview above. It will ask you for your domain name, legitimate email address and then go ahead and copy the templates over, replace placeholder domain names and email address, generate passwords for postgres etc. and then run the production stack. Remember you need to have ufw, rpl, make and pwgen installed before running this command:

```
make configure
```



## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`



## 4.1 Postgres and PostGIS

PostgreSQL is a powerful, open source object-relational database management system (ORDBMS) that uses and extends the SQL language combined with many features that safely store and scale the most complicated data workloads.[1][2]

PostGIS is a spatial extension to for PostgreSQL object-relational database. This allows for GIS (Geographic Information Systems) objects to be stored in the database.[3]

The heart of any stack like the OpenSource GIS Stack would be PostGIS and PostgreSQL. This is because many of the other tools provided in the OSGS stack, for example Goesserver, QGIS Server and Mergin, all rely or can make use of PostGIS as a data storage platform and in some cases as a data analysis platform. The OSGS stack provides PostgreSQL version 13 along with PostGIS which are provided by the the Kartoza PostGIS Docker container. There are a few considerations that were made when deploying this container. One, which is actually a general design consideration of the OSGS stack, is to try and make everything as secure as possible out of the box, rather than leaving security as an afterthought. To fulfill this consideration, what the OSGS stack does, is it spins up a PostgreSQL instance with PostGIS deployed in it and creates a separate user account for connecting to this instance. It also enables SSL encryption by default, as required to the PostgreSQL database. This is so that if you choose to open the port as a public service, firstly, we are not using some well documented default password and username combination and secondly, all the traffic between your client application and the server is encrypted.

**Service name:** db

**Project Website:** [PostgreSQL](#) and [PostGIS](#)

**Project Source Repository:** [postgres](#) / [postgres](#) and [postgis](#) / [postgis](#)

**Project Technical Documentation:** [PostgreSQL Documentation](#) and [PostGIS Documentation](#)

**Docker Repository:** [kartoza/postgis](#)

**Docker Source Repository:** [kartoza](#) / [docker-postgis](#)

### 4.1.1 Configuration

The project includes detailed documentation so this section only contains details relevant to the Open Source GIS Stack configuration.

**Database password:**

Generate a strong password:

```
pwgen 20 1
```

Replace the default docker password for the postgres user with the strong password:

```
rpl "POSTGRES_PASSWORD=docker" "POSTGRES_PASSWORD=<strong password>" .env
```

**Service file configuration:**

Service files entries serve two scenarios:

1. They are needed for opening QGIS projects stored in postgres with PG connection URI because at the project URI you cannot use QGIS authdb. If you prefer to store your projects on the file system, you should rather remove these lines (whole nginx section) since the authentication from pg\_conf/pg\_service.conf can be done more securely by QGIS authdb.
2. Used by your QGIS Server projects to connect to the database once the project is opened from either the file system or the database. You can either specify your password and username in service file or for more advanced configuration you can store user / password credentials in a QGIS authdb file. Refer to the authdb section and in qgis\_conf/qgis-auth.db and the readme in that folder.

On your local machine you should create your own service file with the same service name but connection details that make sense when using the database from your local machine. When you upload your projects into the stack they will connect using the settings from the server hosted service file below assuming you used the same service name.

To carry out the service file configuration, copy, rename then edit the pg\_service file in pg\_config as per the example below (note that we also substitute in the database password created in the steps above).

```
cp pg_conf/pg_service.conf.example \ pg_conf/pg_service.confpassword=docker  
rpl password=<your password> pg_conf/pg_service.conf
```

**Deployment**

```
docker-compose --profile=postgres up -d
```

Note that the default configuration opens the postgresql service to all hosts. This is a potential security hole. If you open the port on the firewall e.g.

```
ufw allow 5432 tcp
```

Then be sure to connect from pg clients like psql or QGIS with SSL enabled so that passwords and data are not transmitted in clear text.



## Validation

Create a local `pg_service.conf` file like the example below and save it in `~/pg_service.conf` or similar as appropriate to your operating system (see <https://www.postgresql.org/docs/12/libpq-pgservice.html> for details on configuration options).

```
[os-gis-stack]
dbname=gis
port=5432
host=<hostname of your server>
user=<your password>
password=docker
```

Now pass the server parameter to `psql` and list the databases as per the example below:

```
[timlinux@fedora ~]$ psql service=os-gis-stack -l
List of databases
```

Name	Owner	Encoding	Collate	Ctype	Access privileges
gis	docker	UTF8	C.UTF-8	C.UTF-8	
postgres	postgres	UTF8	C.UTF-8	C.UTF-8	
template0	postgres	UTF8	C.UTF-8	C.UTF-8	=c/postgres
template1	postgres	UTF8	C.UTF-8	C.UTF-8	=c/postgres

(4 rows)

Test from QGIS is similar:

XXXXXXXXXXXXXXXXXX

Note that there was no need to supply any credentials other than the service file name.

### 4.1.2 Enabling

### 4.1.3 Disabling

### 4.1.4 Accessing the running services

### 4.1.5 Additional Notes

### 4.1.6 References

[1] The PostgreSQL Global Development Group. (n.d.). PostgreSQL: About. PostgreSQL: The World's Most Advanced Open Source Relational Database. Retrieved August 22, 2021, from <https://www.postgresql.org/about/>

[2] The PostgreSQL Global Development Group. (2021, August 12). 1. What Is PostgreSQL? PostgreSQL Documentation. <https://www.postgresql.org/docs/current/intro-what-is.html>

[3] The PostGIS Development Group. (2021, August 20). PostGIS 3.1.4dev Manual. PostGIS - Documentation. <https://postgis.net/docs/manual-3.1/>

## 4.2 GeoServer

GeoServer is a Java-based server, built on [GeoTools](#), that allows users to view and edit geospatial data, using the open standards set forth by the [Open Geospatial Consortium \(OGC\)](#). By implementing the [Web Map Service \(WMS\)](#) standard, GeoServer can create maps in a variety of output formats. GeoServer also conforms to the [Web Feature Service \(WFS\)](#) standard, and [Web Coverage Service \(WCS\)](#) standard which permits the sharing and editing of the data that is used to generate the maps. GeoServer also uses the [Web Map Tile Service](#) standard to split your published maps into tiles for ease of use by web mapping and mobile applications. [OpenLayers](#), a free mapping library, is integrated into GeoServer, making map generation quick and easy. GeoServer is a modular application with additional functionality added via extensions.<sup>[1]</sup>

**Service name:** geoserver

**Project Website:** [GeoServer](#)

**Project Source Repository:** [geoserver / geoserver](#)

**Project Technical Documentation:** [GeoServer Documentation](#)

**Docker Repository:** [kartoza/geoserver](#)

**Docker Source Repository:** [kartoza / docker-geoserver](#)

### 4.2.1 Configuration

```
make configure-geoserver-passwd
```

### 4.2.2 Deployment

```
make deploy-geoserver
```

### 4.2.3 Enabling

```
make enable-geoserver
```

### 4.2.4 Disabling

```
make disable-geoserver
```

### 4.2.5 Starting

```
make start-geoserver
```

## 4.2.6 Logs

```
make geoserver-logs
```

## 4.2.7 Accessing the running services

The services can be accessed on /geoserver/ e.g. <https://localhost/geoserver>.

Look in the .env file for the administrator password.

## 4.2.8 Additional Notes

After configuring, you should remember to set the master password as per [https://docs.geoserver.org/solutions.it/edu/en/security/security\\_overview.html#the-master-password](https://docs.geoserver.org/solutions.it/edu/en/security/security_overview.html#the-master-password) (which is different to the admin password).

## 4.2.9 References

[1] Open Source Geospatial Foundation. (n.d.). About - GeoServer. GeoServer. Retrieved August 21, 2021, from <http://geoserver.org/about/>

## 4.3 Hugo Watcher

Hugo is a static site generator. Hugo builds pages when you create or update your content. Since websites are viewed far more often than they are edited, Hugo is designed to provide an optimal viewing experience for the website's end users and an ideal writing experience for website authors. [1]

The Hugo Watcher service watches for changes in the static content source files of the hugo site and rebuilds the site whenever a source file is changed.

**Service name:** hugo-watcher

**Project Website:** [HUGO](#)

**Project Source Repository:** [gohugoio / hugo](#)

**Project Project Technical Documentation:** [Hugo Documentation](#)

**Docker Repository:** [kartoza/hugo-watcher](#)

**Docker Source Repository:** [kartoza / hugo-watcher](#)

### 4.3.1 Configuration

### 4.3.2 Deployment

### 4.3.3 Enabling

### 4.3.4 Disabling

### 4.3.5 Accessing the running services

### 4.3.6 Additional Notes

### 4.3.7 References

[1] Hugo Authors. (2020, June 3). About Hugo. Hugo. <https://gohugo.io/about/what-is-hugo/>

## 4.4 Docker-osm

OpenStreetMap (OSM) is a digital map database of the world built through crowdsourced volunteered geographic information. The data from OSM is freely available for visualization, query, download, and modification under [open licenses](#). [1] OSM can also be described as a free, editable map of the whole world [2].

The Docker-osm service is a docker compose project to setup an OSM PostGIS database with automatic updates from OSM periodically. The only files you need is a PBF file, geojson (if you intend to restrict data download to a smaller extent than the one specified by the PBF) and run the docker compose project.[3]

**Service name:** osm-mirror

**Project Website:** [OpenStreetMap](#)

**Project Source Repository:** [openstreetmap](#) / [openstreetmap-website](#)

**Project Technical Documentation:** [OpenStreetMap Getting Help](#)

**Docker Repository:** [kartoza/docker-osm](#)

**Docker Source Repository:** [kartoza](#) / [docker-osm](#)

### 4.4.1 Configuration

```
make configure-osm-mirror
```

### 4.4.2 Deployment

```
make deploy-osm-mirror
```

### 4.4.3 Enabling

```
make enable-osm-mirror
```

### 4.4.4 Disabling

```
make ddisable-osm-mirror
```

### 4.4.5 Starting

```
make start-osm-mirror
```

### 4.4.6 Logs

### 4.4.7 Accessing the running services

### 4.4.8 Additional Notes

### 4.4.9 References

[1] Quinn, S., & Dutton, J. A. (n.d.). OpenStreetMap and its use as open data | GEOG 585: Web Mapping. GEOG 585 Open Web Mapping. Retrieved August 30, 2021, from <https://www.e-education.psu.edu/geog585/node/738>

[2] About OpenStreetMap - OpenStreetMap Wiki. (n.d.). OpenStreetMap Wiki. Retrieved August 30, 2021, from [https://wiki.openstreetmap.org/wiki/About\\_OpenStreetMap](https://wiki.openstreetmap.org/wiki/About_OpenStreetMap)

[3] Kartoza. (n.d.). GitHub - kartoza/docker-osm: A docker compose project to setup an OSM Post-GIS database with automatic updates from OSM periodically. GitHub. Retrieved August 30, 2021, from <https://github.com/kartoza/docker-osm#readme>

## 4.5 Lizmap

Lizmap is an open source software designed by 3Liz that allows QGIS® Desktop to create Web map applications. [1]

**Project Website:** [Lizmap](#)

**Project Source Repository:** [3liz / lizmap-web-client](#)

**Project Technical Documentation:** [Lizmap - Documentation](#)

**Docker Repository:** [3liz/lizmap-web-client](#)

**Docker Source Repository:** [3liz / docker-lizmap-web-client](#)

### 4.5.1 Configuration

### 4.5.2 Deployment

### 4.5.3 Enabling

### 4.5.4 Disabling

### 4.5.5 Accessing the running services

### 4.5.6 Additional Notes

### 4.5.7 References

[1] 3Liz. (n.d.). Lizmap. Retrieved August 26, 2021, from <https://www.3liz.com/en/lizmap.html>

## 4.6 Mergin Server

Mergin is a web platform for storage and synchronisation of geospatial projects across multiple users and devices (desktop and mobile). The platform is especially useful when you need:

- **Mobile data collection.** If you need to capture location of assets (and their attributes) or update an existing database.
- **Data sharing.** No complicated setup of access by IT admins to get your data to colleagues or clients. Set up permissions and send invites with few clicks.
- **Offline access.** Work with data with no interruption even without constant Internet connection - sync any changes when you are back online.
- **Collaborative editing.** No more problems dealing with multiple copies of the same dataset in different versions - all changes are automatically consolidated in one place.
- **Audit changes.** Knowing who has changed what and when in a database is often important - Mergin keeps track of the history and allows to go back if needed.
- **No coding required.** Everything can be set up with no knowledge of programming. [1]

**Project Website:** [Mergin](#)

**Project Source Repository:** [lutraconsulting / mergin](#)

**Project Technical Documentation:** [Mergin Help](#)

**Docker Repository:** [lutraconsulting/mergin](#)

**Docker Source Repository:** [lutraconsulting / mergin](#)

### 4.6.1 Configuration

### 4.6.2 Deployment

### 4.6.3 Enabling

### 4.6.4 Disabling

### 4.6.5 Accessing the running services

### 4.6.6 Additional Notes

### 4.6.7 References

[1] Lutra Consulting. (n.d.). GitHub - lutraconsulting/mergin: Store and track changes to your geo-data. GitHub. Retrieved August 26, 2021, from <https://github.com/lutraconsulting/mergin>

## 4.7 Mergin Client

This tool will synchronise a Mergin cloud project into a PostgreSQL project.

There are two modalities in which you can work with Mergin projects:

1. ***mergin-db-sync***: A Mergin project which is synchronised into a PostgreSQL database and supports bidirectional syncing and editing.
2. ***mergin-client* (covered here)**: A folder containing multiple mergin projects (all of the projects shared with a mergin user). These projects are synchronised into the filesystem and published via QGIS Server as web mapping services.

**Service name:** mergin-client

**Project Website:** [Mergin](#)

**Project Source Repository:** [lutraconsulting / mergin-py-client](#)

**Project Technical Documentation:** [Mergin Python Client](#)

**Docker Repository:** [kartoza/mergin-client](#)

**Docker Source Repository:** [kartoza /mergin-client](#)

## 4.7.1 Configuration

## 4.7.2 Deployment

## 4.7.3 Enabling

## 4.7.4 Disabling

## 4.7.5 Accessing the running services

## 4.7.6 Additional Notes

For field data collection and project synchronisation support see the following:

1. The [Mergin](#) platform for cloud storage of projects
2. The [INPUT](#) mobile data collection platform

You use the Mergin client to clone one or more Mergin projects to the host running docker-compose and these projects are made available through QGIS Server.

One critical note is that the Project directory and the Project File names must be the same, otherwise QGIS Server will not recognize the project as being valid. For example:

- Valid: FooProject/FooProject.qgz
- Not Valid: FooProject/BarProject.qgz

Once published in this way, valid projects will be accessible from any OGC compliant client (e.g. QGIS Desktop, OpenLayers, Leaflet) using the following URL Scheme:

`https://yourhost.com/ogc/yourproject`

For example, here is one we published for a client (domain name changed):

`https://example.org/ogc/Elevation`

You can read more about the mergin-client at the separate git repo [here](#).

## 4.8 Mergin-db-sync

There are two modalities in which you can work with Mergin projects:

1. ***mergin-db-sync* (covered here):** A Mergin project which is synchronised into a PostgreSQL database and supports bidirectional syncing and editing.
2. ***mergin-client*:** A folder containing multiple mergin projects (all of the projects shared with a mergin user). These projects are synchronised into the filesystem and published via QGIS Server as web mapping services.

**Service name:** mergin-sync

**Project Website:** [Mergin](#)

**Project Source Repository:** [lutraconsulting / mergin-db-sync](#)

**Project Technical Documentation:** [DB Sync Script](#)

**Docker Repository:** [lutraconsulting/mergin-db-sync](#)

**Docker Source Repository:** [lutraconsulting / mergin-db-sync](#)



## 4.8.1 Configuration

## 4.8.2 Deployment

## 4.8.3 Enabling

## 4.8.4 Disabling

## 4.8.5 Accessing the running services

## 4.8.6 Additional Notes

In the `.env` file you should specify these options:

- **MERGIN\_USER:** This is the user account that will be used to log in and pull/push updates to the Mergin project.
- **MERGIN\_PASSWORD:** This is the user password for the above account.
- **MERGIN\_PROJECT\_NAME:** Specified in the form of `user/project` this is the Mergin project that will be synchronised into the database.
- **MERGIN\_SYNC\_FILE:** This is the name of a GeoPackage `yourgeopackage.gpkg` in the Mergin project whose schema will be replicated into a PostGIS schema as described below.
- **DB\_SCHEMA\_MODIFIED:** This is a PostgreSQL schema (schemas can be thought of as ‘folders’ within your database within which tables are found) that will contain the synchronised data form mergin. The content of the tables are editable via INSERT/UPDATE/DELETE operations, bit the structure of these tables (via ALTER commands) should not be attempted. Note that the replication is bidirectional, so changes made in the database are synchronised to all mergin clients and changes made in the distributed clients will make their way back into your database.
- **DB\_SCHEMA\_BASE:** This is a ‘hands off’ copy of the content in the `MERGIN_SYNC_FILE` that is stored in PostgreSQL to act as a reference when mergin calculates the changeset between the `MODIFIED` schema content and the remote copies of the data. **DO NOT USE THIS** and definately **DO NOT EDIT THIS**..
- **MERGIN\_URL:** This is the public server where your mergin project is hosted. By default it would be “`https://public.cloudmergin.com`” unless you are self hosting the Mergin backend, or using an alternative hosted instance.

Note that in the docker-compose file, the assumption is made that the database being used for Mergin syncing is called ‘gis’ and the hostname (in the private docker network) is called ‘db’. The username and password are taken from the following keys in the `.env` file:

- `POSTGRES_USER`
- `POSTGRES_PASSWORD`

## 4.9 Nginx

Nginx is a lightweight web server acting as a proxy in front of QGIS server, and as a server for the static HTML content.

**Service name:** `nginx`

**Project Website:** [NGINX](#)

**Project Source Repository:** [nginx / nginx](#)

**Project Technical Documentation:** [nginx documentation](#)

**Docker Repository:** [nginx](#)

**Docker Source Repository:** [nginxinc / docker-nginx](#)

## 4.9.1 Configuration

## 4.9.2 Deployment

## 4.9.3 Enabling

## 4.9.4 Disabling

## 4.9.5 Accessing the running services

## 4.9.6 Additional Notes

# 4.10 PostgREST

PostgREST is a standalone web server that turns your PostgreSQL database directly into a RESTful API. The structural constraints and permissions in the database determine the API endpoints and operations. [1] The PostgREST service in the OSGS platform is used for pushing readings from IoT devices to our database.

**Service name:** `postgrest`

**Project Website:** [PostgREST Documentation](#)

**Project Source Repository:** [PostgREST / postgrest](#)

**Project Technical Documentation:** [PostgREST Documentation](#)

**Docker Repository:** [postgrest/postgrest](#)

**Docker Source Repository:** [PostgREST / postgrest](#)

## 4.10.1 Configuration

## 4.10.2 Deployment

## 4.10.3 Enabling

## 4.10.4 Disabling

## 4.10.5 Accessing the running services

## 4.10.6 Additional Notes

## 4.10.7 References

[1] Nelson, J., & Chavez, S. (n.d.). PostgREST Documentation — PostgREST 8.0.0 documentation. PostgREST Documentation. Retrieved August 26, 2021, from <https://postgrest.org/en/v8.0/>

## 4.11 QGIS Server

With the QGIS Server service you can publish one or more QGIS projects including: 1. Projects stored in-database in PostgreSQL 2. Projects stored in the file system

**Project Website:** [QGIS.org](http://QGIS.org)

**Project Source Repository:** [QGIS on GitHub](#)

**Project Project Technical Documentation:** [QGIS on GitHub](#)

**Docker Repository:**

**Docker Source Repository:** [QGIS Server Docker Image](#) from OpenQuake.

### 4.11.1 Configuration

### 4.11.2 Deployment

### 4.11.3 Enabling

### 4.11.4 Disabling

### 4.11.5 Accessing the running services

Every project you publish will be available at `/ogc/project_name` which makes it very simple to discover where the projects are deployed on the server.

### 4.11.6 Additional Notes

For the QGIS Server, we have chosen the OpenQuake build of QGIS Server because it has a few interesting characteristics. One, is that you can deploy QGIS Server side extensions easily with it and two, is that it supports things like the QGIS Authentication System. The QGIS Authentication System is an authentication database that provides more advanced security options, provides Pg service support, and provides some special features for url rerouting so that your project paths are hidden away from the user (which is both a security and a convenience concern).

The OpenQuake QGIS Server is deployed as a QGIS Server instance. The OSGS stack also provides a couple of sample plugins like a demonstrater plugin and a plugin for handling atlas reports. The demonstrater plugin is a modified version of the Git feature info handler and will return some html back and some nicely formatted table. The plugin for handling atlas reports, written by Lizmap and (?) as a QGIS propriety extension to the WMS service protocol, allows you to request a specific page of an atlas, a QGIS composed atlas, as the report. This is pretty handy if you, for example, click on a feature and you want to get then from an atlas report the one page that that feature is covered in the atlas.

Another feature that Docker provides for applications such as QGIS Server is the ability to horizontally scale them. Our platform has some key configuration examples showing you how you can, for example, scale up the QGIS Server instance to have ten concurrently running instances. This is useful for handling increased or high load on the server. It will do like a ram rob(?) and server handler, so that as the requests come in, it will pass each successive request over to the next running instance, and those requests will be handled by that instance, passed back and then that instance will stand by and wait for the next request to come in.

The QGIS Server works in orchestration with many of the other containers, including the PostGIS container. It also works pretty well in conjunction with the SCP (secure copy) container which allows the users of the OSGS architecture to easily move data from their local machine onto the server, either manually by copying and pasting files using an application such as Onescp or using built into Linux file browsers. For example, if you are one the GNOME desktop it has built into SFTP support.

### 4.11.7 Further Reading

You should read the [QGIS Server documentation](#) on QGIS.org. It is well written and covers a lot of background explanation which is not provided here. Also you should familiarise yourself with the [Environment Variables](#).

Alessandro Passoti has made a number of great resources available for QGIS Server. See his [workshop slide deck](#) and his [server side plugin examples](#), and [more examples here](#).

### 4.11.8 QGIS Server Atlas Print Plugin

See the [project documentation](#) for supported request parameters for QGIS Atlas prints.

### 4.11.9 QGIS Server Scaling

If your server has the needed resources, you can dramatically improve response times for concurrent QGIS server requests by scaling the QGIS server:

```
docker-compose --profile=qgis-server up -d --scale qgis-server=10 --remove-orphans
```

To take advantage of this, the `locations/upstreams/qgis-server.conf` should have one server reference per instance e.g.

```
upstream qgis-fcgi {  
    # When not using 'host' network these must reflect the number  
    # of containers spawned by docker-compose and must also have  
    # names generated by it (including the name of the stack)  
    server osgisstack_qgis-server_1:9993;  
    server osgisstack_qgis-server_2:9993;  
    server osgisstack_qgis-server_3:9993;  
    server osgisstack_qgis-server_4:9993;  
    server osgisstack_qgis-server_5:9993;  
    server osgisstack_qgis-server_6:9993;  
    server osgisstack_qgis-server_7:9993;  
    server osgisstack_qgis-server_8:9993;  
    server osgisstack_qgis-server_9:9993;  
    server osgisstack_qgis-server_10:9993;  
}
```

Then restart Nginx too:

```
docker-compose --profile=production restart nginx
```

Note that if you do an Nginx up it may bring down your scaled QGIS containers so take care.

Finally check the logs of Nginx to make sure things are running right:

```
docker-compose --profile=production logs nginx
```

## 4.12 SCP

The SCP (secure copy) containers have been arranged so that there are some standard containers out of the box. Each container has its data stored in its own docker volume as well. The data is somewhat isolated and there are containers for QGIS projects, fonts, SVGs that your QGIS projects might reference, general file sharing, uploading data to ODM, etc. The SCP service is designed to only support connections with SSH public-private key encryption and password based authentication. The way that you provision users into it is that you copy the SSH public key into a file in the configuration folder for SCP and then that user will be allowed to make the connection to whichever SCP share that you have created for them. The SCP container can be used to copy a QGIS project file from your desktop up to the server with all the QGIS resources that it needs such as shapefiles. The QGIS Server instance can then be used to access the project from the OGC web services.

**Project Website:**

**Project Source Repository:**

**Project Project Technical Documentation:**

**Docker Repository:**

**Docker Source Repository:**

### 4.12.1 Configuration

### 4.12.2 Deployment

### 4.12.3 Enabling

### 4.12.4 Disabling

### 4.12.5 Accessing the running services

### 4.12.6 Additional Notes

### 4.12.7 SCP File Drop Shares

This is a container intended for users to upload files for publication on the server. It runs on port 2222 so we need to expose that through the firewall:

```
sudo ufw allow 2222
```

You can add your public keys from the host e.g.

```
cat ~/.ssh/authorized_keys > scp_conf/gis_projects
```

Or copy them in by other means. Each file you create in scp\_conf will be a user name when the scp container runs, with it's own directory in the storage volume, unless an explicit storage volume has been pre-defined (see list of these below). Each file should contain a list of public keys. If you add a key at some point, or a new user file, you may need to restart the container:

```
docker-compose profile=scp restart
```

The following scp shares are made for the various purposes listed below. You need to follow the same pattern of creating a config file for each. These shares each have a dedicated volume associated with it which is also mounted into the associated server container.

---

- **User:** geoserver\_data
  - **Named Volume:** scp\_geoserver\_data
  - **Volume Mounted To:** scp, geoserver
  - **Notes:** Copy vector and raster datasets here for publishing in GeoServer.
  - **Example Use:** sftp://geoserver\_data@<hostname>:2222/home/geoserver\_data
- 

- **User:** qgis\_projects
  - **Named Volume:** scp\_qgis\_projects
  - **Volume Mounted To:** scp, qgis-server
  - **Notes:** Copy QGIS projects and data here for publishing with QGIS Server. See notes on directory layout below.
  - **Example Use:** sftp://qgis\_projects@<hostname>:2222/home/qgis\_projects
- 

- **User:** qgis\_svgs
  - **Named Volume:** scp\_qgis\_svgs
  - **Volume Mounted To:** scp, qgis-server
  - **Notes:** Embed SVGs in styles by preference in QGIS. Use this drop if you have no way to use embedded SVGs.
  - **Example Use:** ``sftp://qgis\_svgs@:2222/home/qgis\_svgs`
- 

- **User:** qgis\_fonts
  - **Named Volume:** scp\_qgis\_fonts
  - **Volume Mounted To:** scp, qgis-server
  - **Notes:** Copy fonts directly into the root folder.
  - **Example Use:** sftp://qgis\_fonts@<hostname>:2222/home/qgis\_fonts
- 

- **User:** hugo\_data
  - **Named Volume:** scp\_hugo\_data
  - **Volume Mounted To:** scp, hugo\*
  - **Notes:** Upload markdown files for static site generation with Hugo.
  - **Example Use:** sftp://hugo\_data@<hostname>:2222/home/hugo\_data
- 

- **User:** odm\_data
  - **Named Volume:** scp\_odm\_data
-

- **Volume Mounted To:** scp, odm \*
- **Notes:** Upload imagery data for processing with ODM
- **Example Use:** sftp://odm\_data@<hostname>:2222/home/odm\_data

- 
- **User:** general\_data
  - **Named Volume:** scp\_general\_data
  - **Volume Mounted To:** scp
  - **Notes:** General sharing directory. Later we will publish this under nginx for public downloads. Don't put any sensitive data in here.
  - **Example Use:** sftp://general\_data@<hostname>:2222/home/general\_data
- 

**Note:** Any user connecting to any of these shares will be able to see all other files from all other users. They will only have write access to the folder they are connecting to, for all other shares their access will be read only. If you want to further partition the access to files you can create multiple scp services, each with one of the mount points listed above. In so doing users would not be able to see the other mount points listed above.

#### 4.12.8 Directory layout for the QGIS projects folder

When adding projects to the qgis\_projects folder, you need to follow this convention strictly for the projects to be recognised by QGIS Server:

qgis\_projects/<project\_name>/<project\_name>.qgs

For example:

qgis\_projects/terrain/terrain.qgs

There is a convenience Make target that will copy your .ssh/authorized\_keys file contents into each of the scp\_config user files listed in the table above.

```
make setup-scp
```

#### 4.12.9 Starting the container

```
docker-compose --profile=scp up -d scp
```

Example copying of data into the container from the command line:

```
scp -P 2222 sample-document.txt localhost:/data//gis_projects/gis_projects/gis_projects
```

In Nautilus (file manager in Linux Gnome Desktop) you can test by connecting

```
sftp://<hostname>:2222/data/gis_projects
```

into the red highlighted box below:

```
XXXXXXXXXXXXXXXXXXXXXXX
```

After that open a second window and you can drag and drop files too and from the folder. Windows users can use the free WinSCP application to copy files to the server.

#### 4.12.10 FAQ

**Q:** When connecting I get “Host key validation failure” or similar **A:** Remove the entry for the server in your `~/.ssh/known_hosts`



## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`



## UTILITIES

## 6.1 make ps

### 6.1.1 Synopsis

Displays the actively running services in the OSGS platform.

### 6.1.2 Usage

**Make Target:** ps

**Arguments:** none

**Example usage:** make ps

**Example output:**

-----  
Current status  
-----

Name	Ports	Command	State
osgisstack_db_1		/bin/sh -c /scripts/docker ...	Up (healthy)
osgisstack_hugo-watcher_1		python3 /hugo_watcher.py	Up
osgisstack_nginx_1		/docker-entrypoint.sh /bin ...	Up
osgisstack_scp_1		/bin/sh -c /run.sh	Up

### 6.1.3 Notes

Make ps only shows active containers.

## 6.2 make kill

### 6.2.1 Synopsis

Kills all actively running services in the OSGS platform.

### 6.2.2 Usage

**Make Target:** kill

**Arguments:** none

**Example usage:** make kill

**Example output:**

```
-----  
Killing all containers  
-----  
Killing osgisstack_db_1          ... done  
Killing osgisstack_nginx_1      ... done  
Killing osgisstack_hugo-watcher_1 ... done  
Killing osgisstack_scp_1        ... done  
    0.0.0.0:2222->22/tcp, :::2222->22/tcp
```

### 6.2.3 Notes

Make kill only shows active containers from the OSGS platform, other containers should continue unaffected.

## 6.3 make nuke

### 6.3.1 Synopsis

Kills all actively running services in the OSGS platform, deletes all configurations, deletes all docker volumes, removes letsencrypt certificate. Generally has the effect of resetting your system to the state it was at when you performed the initial git checkout.

### 6.3.2 Usage

**Make Target:** nuke

**Arguments:** none

**Example usage:** make nuke

**Example output:**

```
-----
Disabling services
This command will delete all your configuration and data permanently.
Are you sure? [y/N] y
Please type CONFIRM to proceed CONFIRM
-----
Nuking Everything!
-----
Going to remove osgisstack_db_1, osgisstack_nginx_1, osgisstack_hugo-watcher_1, ↵
↵osgisstack_scp_1
Removing osgisstack_db_1          ... done
Removing osgisstack_nginx_1      ... done
Removing osgisstack_hugo-watcher_1 ... done
Removing osgisstack_scp_1        ... done

[sudo] password for timlinux:
make[1]: Entering directory '/home/timlinux/dev/docker/OpenSource-GIS-Stack'

-----
Reset site configuration to default values
This will replace any local configuration changes you have made
-----
Are you sure you want to continue? [y/N] y
make[1]: Leaving directory '/home/timlinux/dev/docker/OpenSource-GIS-Stack'
make[1]: Entering directory '/home/timlinux/dev/docker/OpenSource-GIS-Stack'

-----
Disabling services
This will remove any symlinks in conf/nginx_conf/locations and conf/nginx_conf/upstreams
effectively disabling all services exposed by nginx
-----
Are you sure? [y/N] y
make[1]: Leaving directory '/home/timlinux/dev/docker/OpenSource-GIS-Stack'
```

### 6.3.3 Notes

Again we warn you here that make nuke is extremely destructive and should only be used if you want to completely reset your installation!

## 6.4 make docs

### 6.4.1 Synopsis

Builds the system documentation.

### 6.4.2 Usage

**Make Target:** docs

**Arguments:** none

**Example usage:** make docs

**Example output:**

**See also:** *building docs*

### 6.4.3 Notes

None

## 6.5 Building Documentation

### 6.5.1 HTML Docs

```
sudo dnf install sphinx
pip3 install --upgrade myst-parser
pip install sphinx-sizzle-theme
```

See <https://sphinx-themes.org/sample-sites/sphinx-sizzle-theme/> for theme specific info.

See <https://www.sphinx-doc.org/en/master/usage/markdown.html> for Markdown in Sphinx support notes (and the docs at <https://myst-parser.readthedocs.io/en/latest/syntax/optional.html>). Especially, note the admonitions docs which are used to make little alert etc boxes in the docs: <https://myst-parser.readthedocs.io/en/latest/syntax/optional.html#html-admonitions>

## 6.5.2 PDF Docs

On fedora I just install the huge, full texlive package:

```
sudo dnf install texlive-scheme-full
```

Then build:

```
make latexpdf
```

Sometimes you need to run it a second time if it is a fresh build.

After the build is done, the PDF will be at:

```
docs/build/latex/osgs.pdf
```





## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`



## **WORKFLOWS**

### **8.1 PostgreSQL Workflows**

There are a number of different workflows we will explore here:

- Connecting to the PostgreSQL database from your QGIS Desktop application
- Connecting to the PostgreSQL database from QGIS Server
- Connecting to the PostgreSQL database from GeoServer
- Connecting from other applications

Before we dive into the details here, let us quickly examine some basic use cases:

- You collect data in the field using Input and use the mergin-db-sync workflow to synchronised field collected data into the database, then publish maps for this data in QGIS Server.
- You want to publish a layer in GeoServer, so you connect to the database from your desktop, drag and drop a local file system layer into the database using QGIS, then publish the layer from GeoServer.
- You want to create a project in QGIS desktop that uses PostgreSQL for data storage, then publish that data as a QGIS Project.

There are many other workflows that the Open GIS Stack supports, it is really up to your imagination! What is key to understand are the mechanisms that you can use to connect to the database in different contexts. And so we will focus this section on the different connection modalities.

#### **8.1.1 Direct connection**

Direct connection to the server over the standard PostgreSQL port exposed to the public internet is probably the least secure but most convenient approach. We do not recommend this approach, and if you do follow it be sure to use the option to force remote clients to connect using SSL (and expect a performance penalty in the process).

Note also that connecting to a PostgreSQL database using QGIS from a remote host may often be slow and irritating to use on a daily basis, especially if your internet connection is not very fast and your datasets are large.

By default the docker-compose.yaml

Publishing with QGIS Server

The workflows described in this section apply equally to any database hosted

Further reading for understanding authentication with PostgreSQL using cert based authentication here:

<https://joelonsql.com/2013/04/27/securing-postgresql-using-hostssl-cert-clientcert1/>



## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`



## OSM MIRROR

The OSM mirror uses the kartoza/docker-osm tool to create an in-database mirror of a designated geographical area in the designated postgres database schema (set to: osm). The OSM mirror tool is described in the project README here:

<https://github.com/kartoza/docker-osm>

To deploy the docker-mirror you need to follow the steps descibed below. First a process overview:

1. Create the PBF feature container passing it a URL to a PBF file
2. Create a clip file that will be used to constrain any retrieved / imported data to a specific geographic area.
3. Tweak the mappings.yml file (advanced users)
4. Run the docker-osm service
5. Optionally include these data in published services via QGIS projects, GeoServer etc.

### 10.1 PBF Container

During the `make configure` process, the script will ask for the URL to an OSM .pbf file e.g.:

```
-----  
Fetching pbf if not cached and then copying to settings dir  
-----  
URL For Country PBF File: https://download.geofabrik.de/central-america-latest.osm.pbf
```

You can enter any valid URL for an OSM .PBF file at this point. A docker container will be built that fetches the PBD and stores it on the host file system under `osm_config`.

### 10.2 Clip Area

Create the clip area to constrain the geographical region that data will be harvested for. For best performance, a simple rectangle is best, but any complex polygon can be used. The clip area must be saved as `osm_config/clip.geojson`. The format for the clip area must be GeoJSON. You can easily create this using QGIS.

## 10.3 Mappings

For advanced users, you can tweak the `osm_config/mapping.yml`

You can see how the `imposm3` mapping syntax works here:

<https://imposm.org/docs/imposm3/latest/mapping.html>

Note that you cannot alter the mappings after the service is running without clearing the database and restarting the import.

## 10.4 Run services

To run the OSM services do:

```
docker-compose --profile=osm up -d
```

## 10.5 Publishing with GeoServer

You can publish the data in the `osm` schema using GeoServer or by publishing a QGIS project that references the data layers in the OSM schema.

The steps for publishing with GeoServer are quite simple:

1. Log in to GeoServer using the ‘admin’ user and the password in `.env`.
2. Create a new store of type ‘Postgis’ and configure it as per the screenshot below, replacing the password with the Postgres password stored in `.env`:

The screenshot shows the GeoServer web interface. The top bar indicates 'Logged in as admin.' with a 'Logout' button. The left sidebar contains a navigation menu with categories: About & Status, Data, Services, and Settings. The main content area is titled 'Edit Vector Data Source' and contains the following configuration fields:

- PostGIS Database**: PostGIS Database
- Basic Store Info**:
  - Workspace \*: SaintLucia
  - Data Source Name \*: OSM Mirror
  - Description: Mirror of OSM data for St Lucia
  - ☒ Enabled
- Connection Parameters**:
  - host \*: db
  - port \*: 5432
  - database: gis
  - schema: osm
  - user \*: docker
  - passwd: [redacted]
  - Namespace \*: saintlucia
  - ☒ Expose primary keys
  - max connections: [empty]

At the bottom of the form are three buttons: 'Save', 'Apply', and 'Cancel'.

Also, be sure to scroll down and set SSL mode to Required:



### Callback factory

☒ Loose bbox

☒ Estimated extends

### SSL mode

☐ preparedStatements

### Max open prepared statements

☒ encode functions

☒ Support on the fly geometry simplification

3. Register one or more layers from that store as per the image below:

**New Layer**

Add a new layer

Add layer from: **SaintLucia:OSM Mirror**

You can create a new Feature type by manually configuring the attribute names and types. [Create new feature type...](#)

On databases you can also create a new feature type by configuring a native SQL statement. [Configure new SQL view...](#)

Here is a list of resources contained in the store 'OSM Mirror'. Click on the layer you wish to configure

Results 0 to 0 (out of 0 items)

Published	Layer name	Action
✓	osm_admin	Publish again
✓	osm_buildings	Publish again
✓	osm_places	Publish again
	osm_aerodrome_label_point	Publish
	osm_aeroway_linestring	Publish
	osm_aeroway_point	Publish
	osm_aeroway_polygon	<b>Publish</b>
	osm_bay	Publish
	osm_border_disp_relation	Publish
	osm_building_polygon	Publish
	osm_building_relation	Publish
	osm_busbar	Publish
	osm_compensator_areas	Publish
	osm_compensator_point	Publish
	osm_converter_areas	Publish
	osm_converter_point	Publish
	osm_healthcare_facilities_node	Publish
	osm_healthcare_facilities_way	Publish
	osm_housenumber_point	Publish

4. Complete the layer details as appropriate and make sure to click the options highlighted in red in the screenshot below:

Security

Settings

Authentication

Passwords

Users, Groups, Roles

Data

Services

WPS security

Monitor

Activity

Reports

Demos

Tools

Add link

Note only FGDC and TC211 metadata links show up in WMS 1.1.1 capabilities

Data links

No data links so far

Add link

Coordinate Reference Systems

Native SRS

EPSG:4326

EPSG:WGS 84...

Declared SRS

EPSG:4326

Find...

EPSG:WGS 84...

SRS handling

Force declared

Bounding Boxes

Native Bounding Box

Min X

Min Y

Max X

Max Y

-61.00221252

13.730535507

-60.88537216

14.022650718

Compute from data

Compute from SRS bounds

Lat/Lon Bounding Box

Min X

Min Y

Max X

Max Y

-61.00221252

13.730535507

-60.88537216

14.022650718

Compute from native bounds

Curved geometries control

Linear geometries can contain circular arcs

Linearization tolerance (useful only if your data contains curved geometries)

Feature Type Details

Property	Type	Nullable	Min/Max Occurrences
id	Integer	false	1/1

Save

Apply


Cancel

1. Connect to the GeoServer from a client e.g. QGIS using WFS or WMS using the scheme:

<https://example.org/geoserver/SaintLucia/wfs>

or

<https://example.org/geoserver/SaintLucia/wms>



## 10.6 Publishing with QGIS Server

The workflows described in the section on working with the PostgreSQL database below are basically all you need to know, so we don't repeat that here, other than to remind you that the OSM mirrored data is by default stored in a schema called 'osm'.

## 10.7 OSM Attribution

Note that whenever you publish a map containing OSM data, be careful to adhere to the license and credit the OSM Project as per:

<https://www.openstreetmap.org/copyright>



## INDICES AND TABLES

- `genindex`
- `modindex`
- `search`